

P8106_group2recovery_secondaryanalysis

Yimin Chen (yc4195), Yang Yi (yy3307), Qingyue Zhuo (qz2493)

Contents

Import and data manipulation	1
Data visualization	3
Correlation plot	3
Feature plot	3
Partition plot	4
Model training	5
Logistic regression and its cousins	5
GLM	5
Penalized logistic regression	5
GAM	6
MARS	7
test data performance for Logistic regression and its cousins	10
Discriminant Analysis	11
LDA	11
QDA	11
Naive Bayes (NB)	11
test set performance for Discriminant Analysis	13
classification tree models	14
rpart	14
ctree	16
test set performance for classification tree models	17
Support Vector Machines	18
test data performance of SVM methods	21
Model comparison	22

Import and data manipulation

```
# Load recovery.RData environment
load("./recovery.Rdata")

dat %>% na.omit()

# dat1 draw a random sample of 2000 participants Uni:3307
set.seed(3307)
```

```

dat1 = dat[sample(1:10000, 2000),]

dat1 =
  dat1[, -1] %>%
  mutate(
    recovery_time = as.factor(
      case_when(recovery_time <= 30 ~ "long", recovery_time > 30 ~ "short")
    ),
    gender = as.factor(gender),
    race = as.factor(race),
    smoking = as.factor(smoking),
    hypertension = as.factor(hypertension),
    diabetes = as.factor(diabetes),
    vaccine = as.factor(vaccine),
    severity = as.factor(severity),
    study = as.factor(
      case_when(study == "A" ~ 1, study == "B" ~ 2, study == "C" ~ 3)
    )
  )

# dat2 draw a random sample of 2000 participants Uni:2493
set.seed(2493)

dat2 = dat[sample(1:10000, 2000),]

dat2 =
  dat2[, -1] %>%
  mutate(
    recovery_time = as.factor(
      case_when(recovery_time <= 30 ~ "long", recovery_time > 30 ~ "short")
    ),
    gender = as.factor(gender),
    race = as.factor(race),
    smoking = as.factor(smoking),
    hypertension = as.factor(hypertension),
    diabetes = as.factor(diabetes),
    vaccine = as.factor(vaccine),
    severity = as.factor(severity),
    study = as.factor(
      case_when(study == "A" ~ 1, study == "B" ~ 2, study == "C" ~ 3)
    )
  )

# Merged dataset with unique observation
covid_dat = rbind(dat1, dat2) %>%
  unique()

covid_dat2 = model.matrix(recovery_time ~ ., covid_dat)[, -1]

# Partition dataset into two parts: training data (70%) and test data (30%)
rowTrain = createDataPartition(y = covid_dat$recovery_time, p = 0.7, list = FALSE)

trainData = covid_dat[rowTrain, ]

```

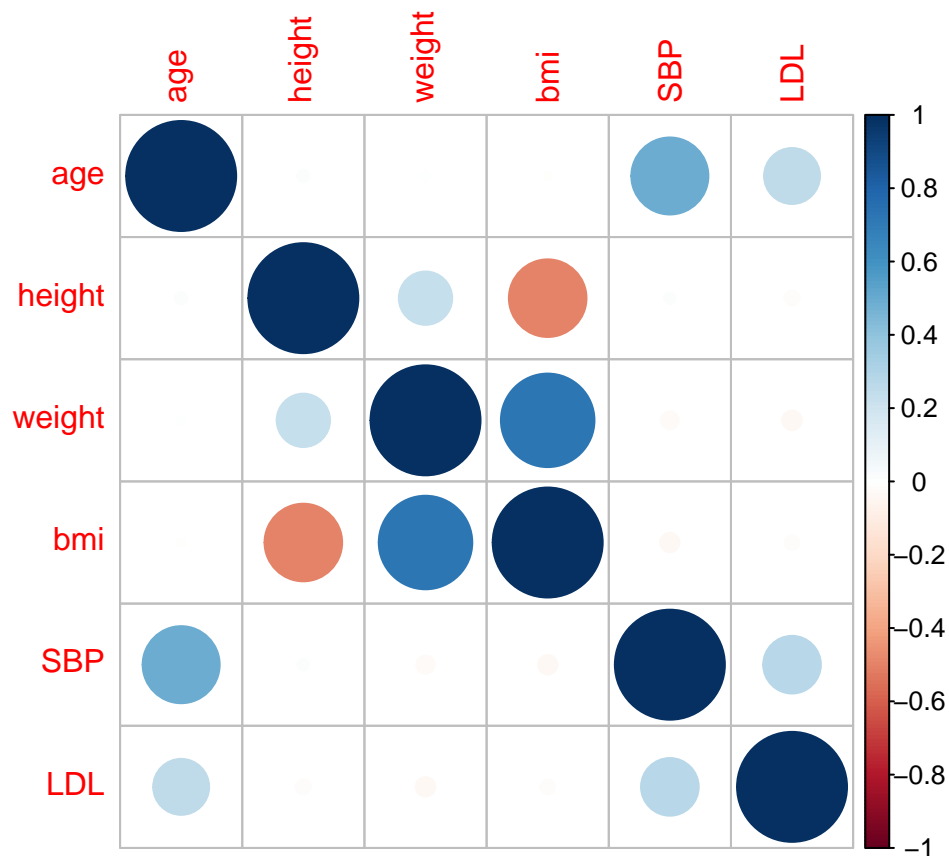
```
testData = covid_dat[-rowTrain, ]

ctrl1 = trainControl(method = "repeatedcv", number = 10, repeats = 5)
ctrl2 = trainControl(method = "cv",
                      classProbs = TRUE,
                      summaryFunction = twoClassSummary)
```

Data visualization

Correlation plot

```
corr_dat = covid_dat[rowTrain,] %>%
  dplyr::select('age', 'height', 'weight', 'bmi', 'SBP', 'LDL')
corrplot(cor(corr_dat), method = "circle", type = "full")
```



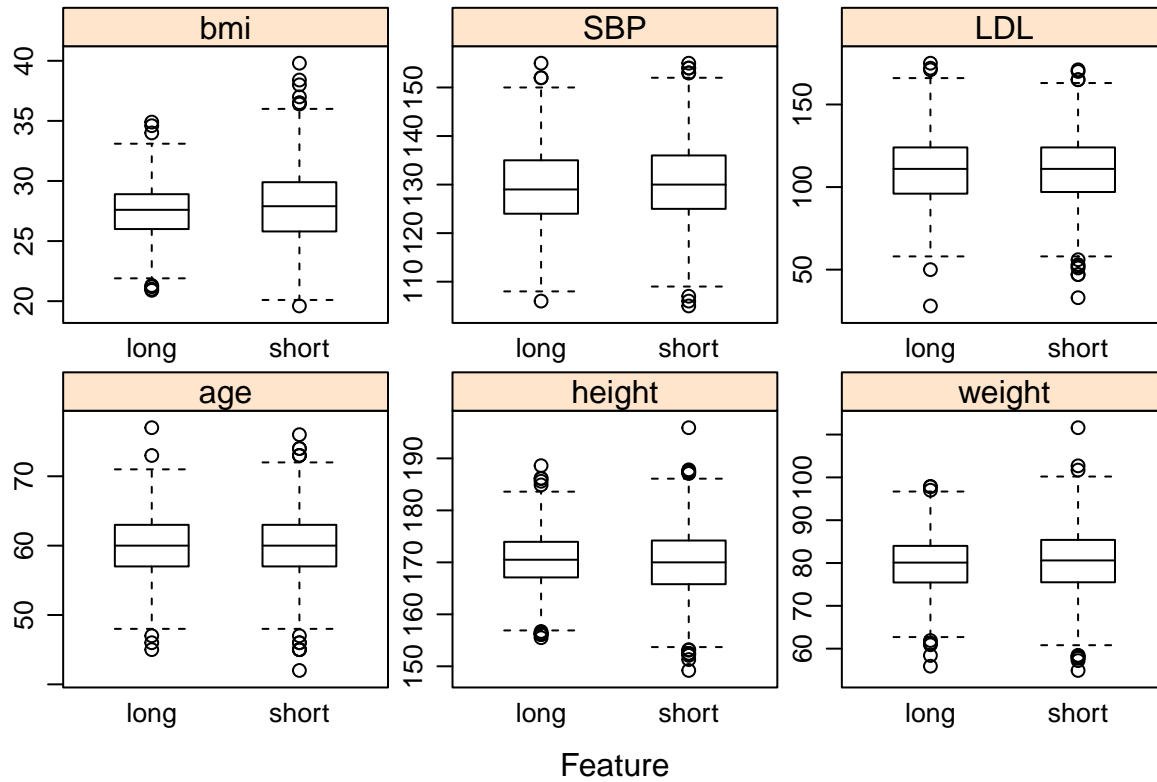
Feature plot

```
vis_trdat = trainData %>%
  dplyr::select('age', 'height', 'weight', 'bmi', 'SBP', 'LDL', 'recovery_time')

theme1 = transparentTheme(trans = .4)
trellis.par.set(theme1)

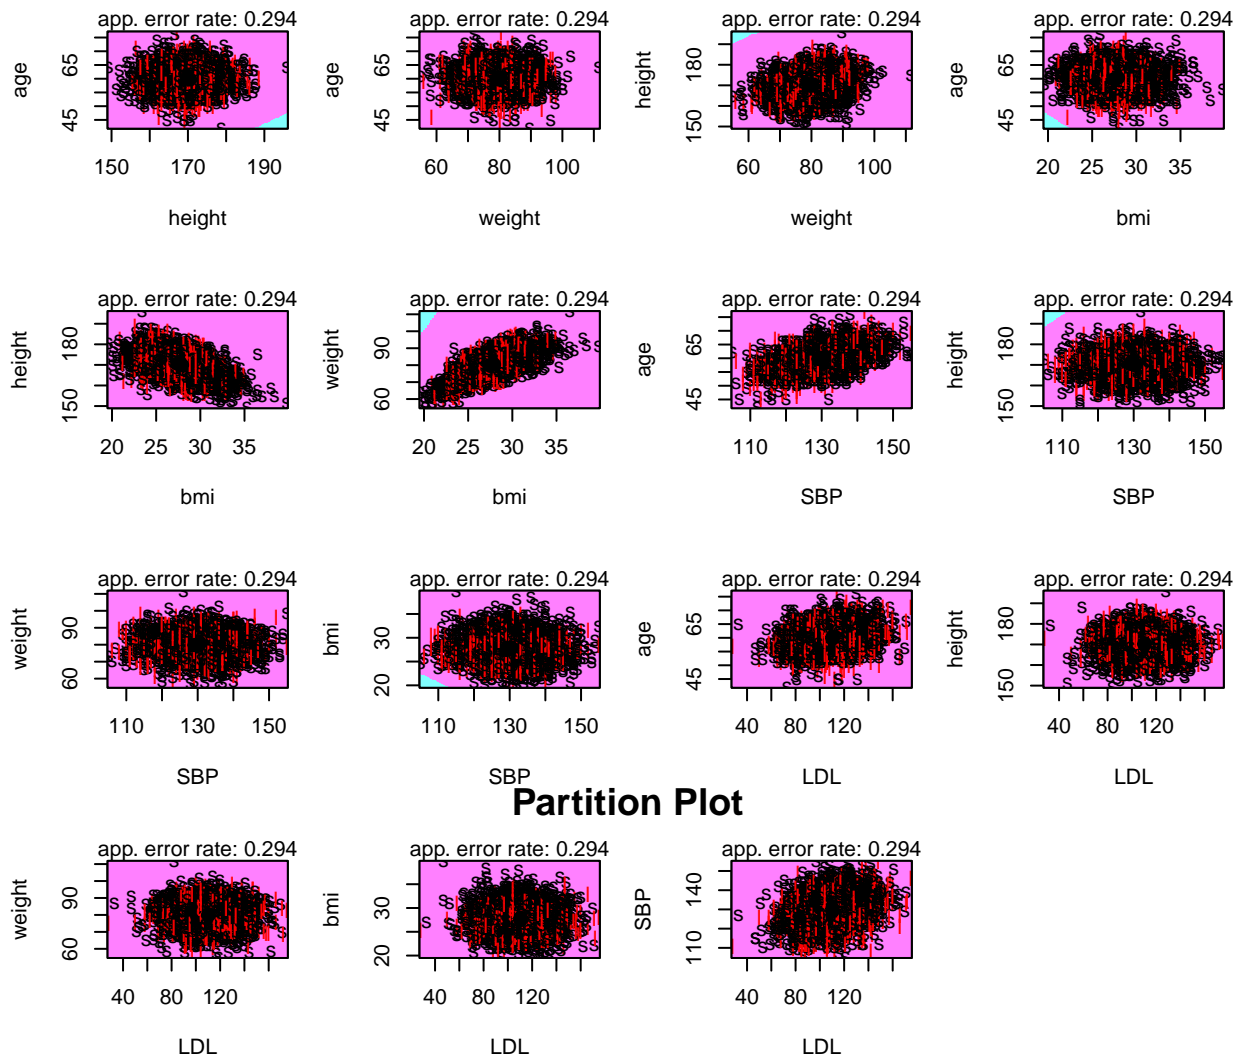
featurePlot(x = vis_trdat[, 1:6],
            y = vis_trdat[, 7],
```

```
scales = list(x = list(relation = "free"),
              y = list(relation = "free")),
plot = "box", pch = "|",
auto.key = list(columns = 2))
```



Partition plot

```
partimat(recovery_time ~ age + height + weight + bmi + SBP + LDL, data = covid_dat, subset = rowTrain, n
```



Model training

Logistic regression and its cousins

GLM

```
set.seed(2)
model.glm <- train(x = covid_dat2[rowTrain,],
  y = covid_dat$recovery_time[rowTrain],
  method = "glm",
  metric = "ROC",
  trControl = ctrl2)
```

Penalized logistic regression

```
glmGrid <- expand.grid(.alpha = seq(0, 1, length = 21),
  .lambda = exp(seq(-8, -1, length = 50)))
set.seed(2)
model.glmn <- train(x = covid_dat2[rowTrain,],
```

```

y = covid_dat$recovery_time[rowTrain],
method = "glmnet",
tuneGrid = glmnGrid,
metric = "ROC",
trControl = ctrl2)


model.glmn$bestTune

##      alpha      lambda
## 1001      1 0.0003354626

myCol<- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
              superpose.line = list(col = myCol))

plot(model.glmn, par.settings = myPar, xTrans = function(x) log(x))

```



GAM


```

set.seed(2)
model.gam <- train(x = covid_dat2[rowTrain,],
                  y = covid_dat$recovery_time[rowTrain],
                  method = "gam",
                  metric = "ROC",
                  trControl = ctrl2)

model.gam$finalModel

```

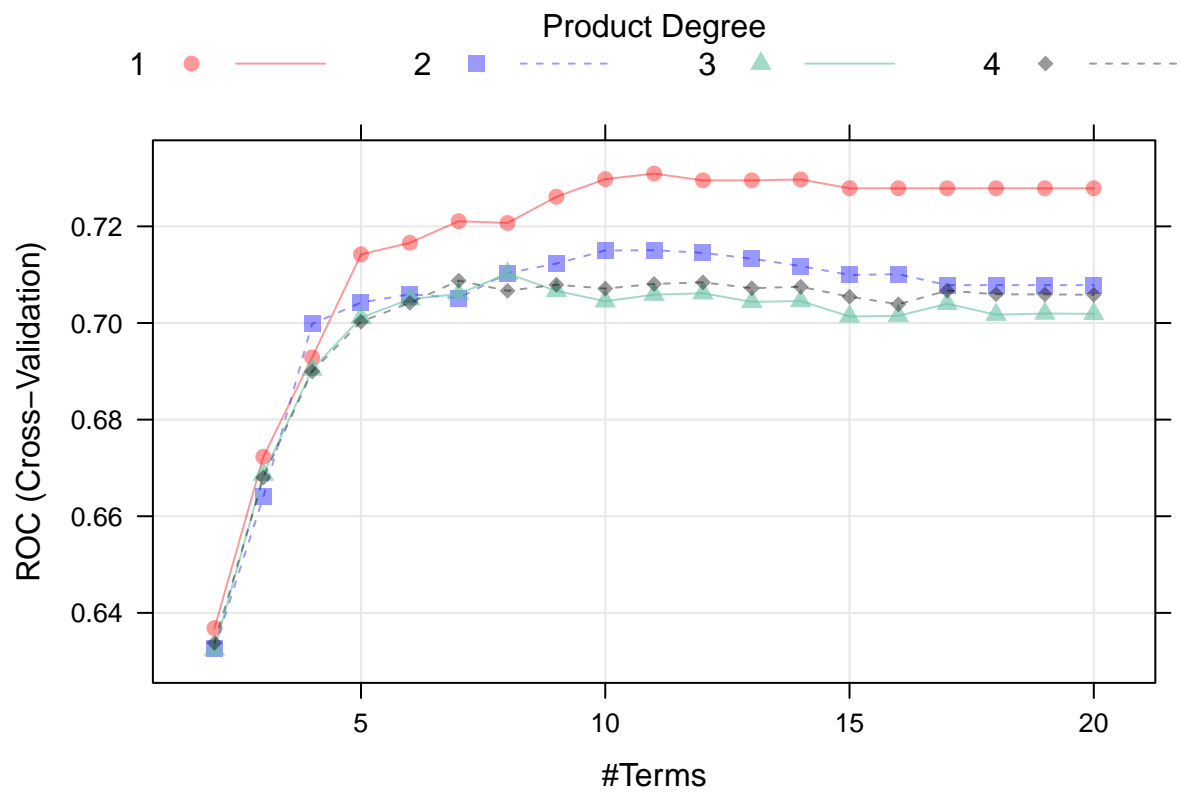
```
##
## Family: binomial
## Link function: logit
##
## Formula:
## .outcome ~ gender1 + race3 + race4 + smoking1 + smoking2 + hypertension1 +
##     diabetes1 + vaccine1 + severity1 + study2 + study3 + s(age) +
##     s(SBP) + s(LDL) + s(bmi) + s(height) + s(weight)
##
## Estimated degrees of freedom:
## 1.00 1.00 1.55 2.79 1.00 2.29 total = 21.63
##
## UBRE score: 0.06815373
plot(model.gam$finalModel, select = 3)
```



MARS

```
set.seed(2)
model.mars <- train(x = covid_dat2[rowTrain,],
                    y = covid_dat$recovery_time[rowTrain],
                    method = "earth",
                    tuneGrid = expand.grid(degree = 1:4,
                                           nprune = 2:20),
                    metric = "ROC",
                    trControl = ctrl12)

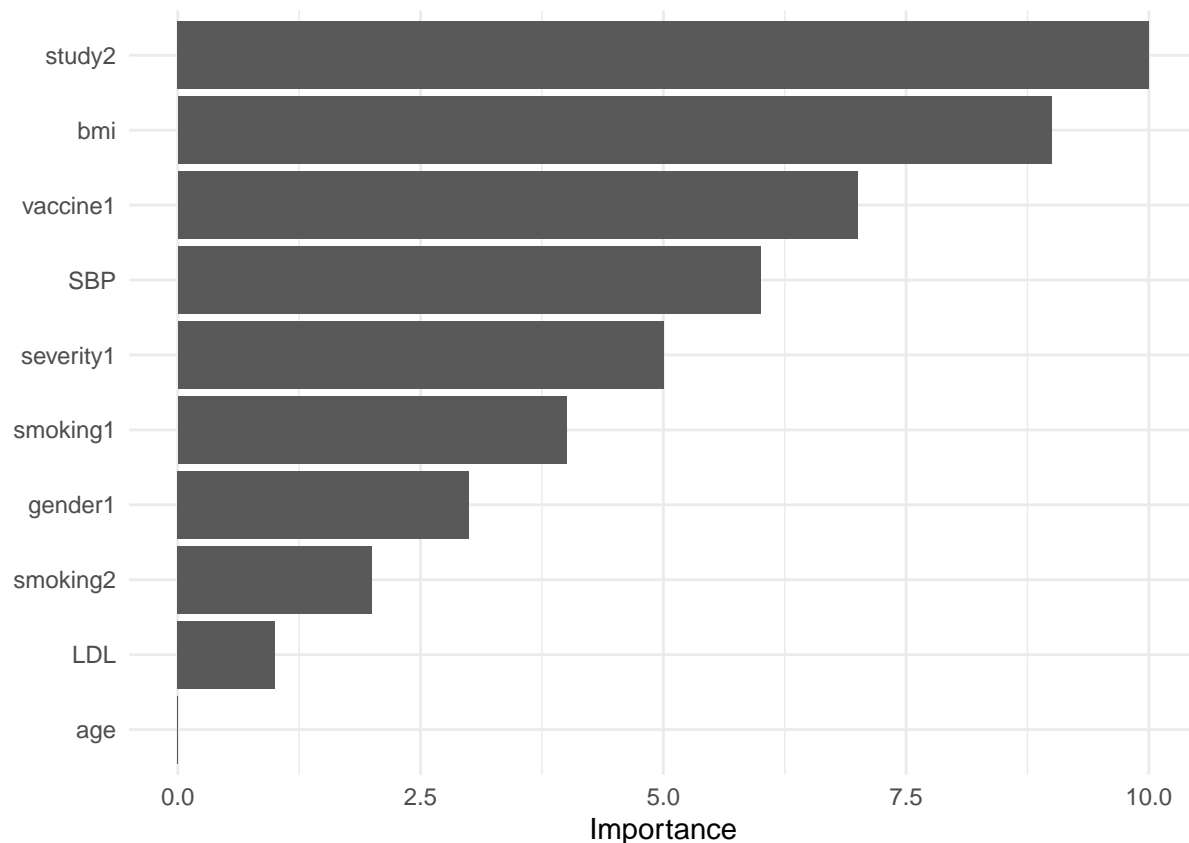
plot(model.mars)
```



```
coef(model.mars$finalModel)
```

```
## (Intercept)      study2 h(28.6-bmi)    vaccine1 h(135-SBP)    severity1
## -0.32524568 -1.35310824  0.51047027 -0.73109733 -0.03262848  0.80307433
##      smoking1      gender1      smoking2 h(LDL-145) h(bmi-23.1)
##  0.43021337 -0.32207625  0.55022116 -0.05342548  0.41456148
```

```
vip(model.mars$finalModel)
```

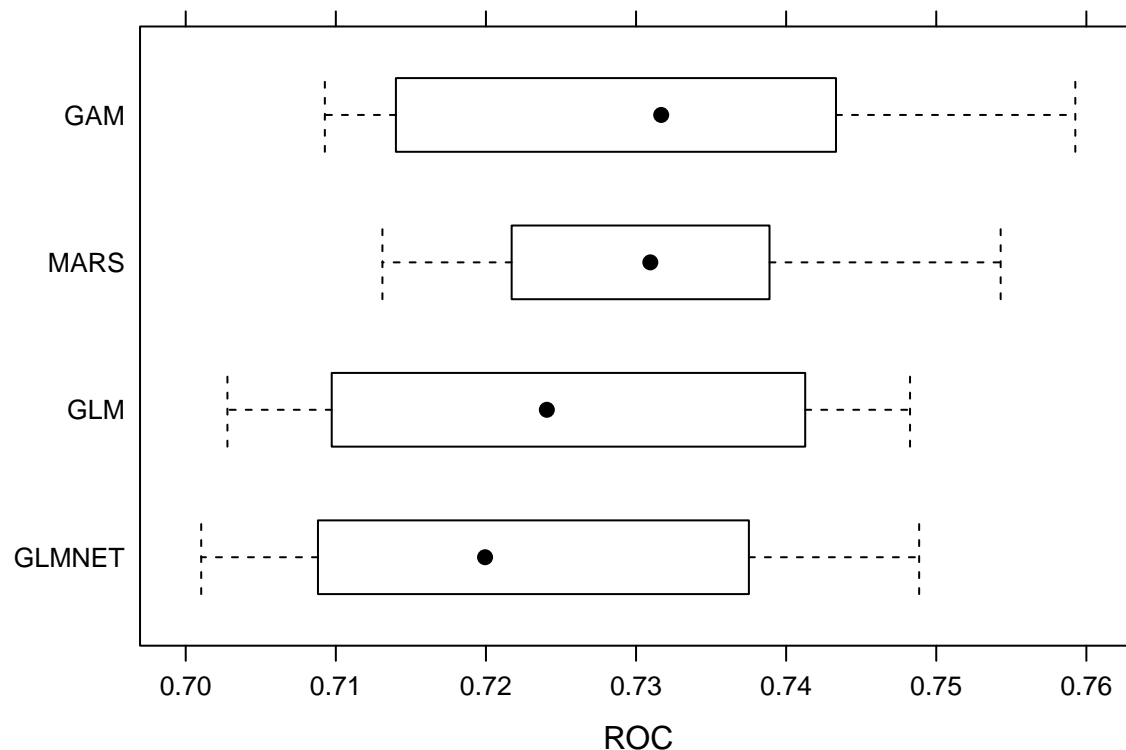



```
res <- resamples(list(GLM = model.glm,
                      GLMNET = model.glmn,
                      GAM = model.gam,
                      MARS = model.mars))
summary(res)
```

```
##
## Call:
## summary.resamples(object = res)
##
## Models: GLM, GLMNET, GAM, MARS
## Number of resamples: 10
##
## ROC
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM    0.7027786 0.7100512 0.7240580 0.7243182 0.7396824 0.7482539    0
## GLMNET 0.7010325 0.7091541 0.7199426 0.7230782 0.7374113 0.7488612    0
## GAM    0.7092686 0.7156054 0.7316766 0.7319728 0.7432725 0.7592621    0
## MARS   0.7131036 0.7220759 0.7309513 0.7309203 0.7385390 0.7542894    0
##
## Sens
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM    0.2162162 0.2702703 0.2789708 0.2814698 0.3074324 0.3378378    0
## GLMNET 0.2027027 0.2466216 0.2837838 0.2693262 0.2969733 0.3243243    0
## GAM    0.2027027 0.2627730 0.3175676 0.3017031 0.3378378 0.3648649    0
## MARS   0.2297297 0.2837838 0.2924843 0.3044428 0.3344595 0.3783784    0
##
```

```
## Spec
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM      0.8531073 0.9039548 0.9154605 0.9119755 0.9324890 0.9438202    0
## GLMNET 0.8644068 0.9053672 0.9180791 0.9187425 0.9382022 0.9548023    0
## GAM      0.8644068 0.8884181 0.9124294 0.9018314 0.9157303 0.9269663    0
## MARS     0.8644068 0.8912429 0.9098584 0.9057862 0.9196106 0.9325843    0
```

```
bwplot(res, metric = "ROC")
```



test data performance for Logistic regression and its cousins


```
glm.pred <- predict(model.glm, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
glmnet.pred <- predict(model.glmnet, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
gam.pred <- predict(model.gam, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
mars.pred <- predict(model.mars, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]

roc.glm <- roc(covid_dat$recovery_time[-rowTrain], glm.pred)
roc.glmnet <- roc(covid_dat$recovery_time[-rowTrain], glmnet.pred)
roc.gam <- roc(covid_dat$recovery_time[-rowTrain], gam.pred)
roc.mars <- roc(covid_dat$recovery_time[-rowTrain], mars.pred)

auc <- c(roc.glm$auc[1], roc.glmnet$auc[1],
         roc.gam$auc[1], roc.mars$auc[1])

modelNames <- c("glm", "glmnet", "gam", "mars")

ggroc(list(roc.glm, roc.glmnet, roc.gam, roc.mars), legacy.axes = TRUE) +
  scale_color_discrete(labels = paste0(modelNames, " (", round(auc, 3), ")"),
                        name = "Models (AUC)") +
  geom_abline(intercept = 0, slope = 1, color = "grey")
```



Models (AUC) — glm (0.691) — glmn (0.691) — gam (0.711) — mars (0.716)

Discriminant Analysis

LDA

```
set.seed(2)

model.lda <- train(x = covid_dat2[rowTrain,],
  y = covid_dat$recovery_time[rowTrain],
  method = "lda",
  metric = "ROC",
  trControl = ctrl2)
```

QDA

```
set.seed(2)

model.qda <- train(x = covid_dat2[rowTrain,],
  y = covid_dat$recovery_time[rowTrain],
  method = "qda",
  metric = "ROC",
  trControl = ctrl2)
```

Naive Bayes (NB)

```
nbGrid <- expand.grid(usekernel = c(FALSE, TRUE),
  fL = 1,
```


```

adjust = seq(.2, 3, by = .2))

set.seed(2)
model.nb <- train(x = covid_dat2[rowTrain,],
  y = covid_dat$recovery_time[rowTrain],
  method = "nb",
  tuneGrid = nbGrid,
  metric = "ROC",
  trControl = ctrl12)

plot(model.nb)

```



```

res <- resamples(list(LDA = model.lda, QDA = model.qda, NB = model.nb))
summary(res)

```

```

##
## Call:
## summary.resamples(object = res)
##
## Models: LDA, QDA, NB
## Number of resamples: 10
##
## ROC
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## LDA 0.6982994 0.7092113 0.7204289 0.7236334 0.7385363 0.7509869    0
## QDA 0.6825470 0.6870515 0.7021897 0.7067165 0.7212097 0.7436988    0
## NB  0.6882730 0.7120171 0.7177622 0.7186405 0.7223240 0.7519739    0
##
## Sens

```

```
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## LDA 0.1891892 0.26013514 0.27702703 0.26525361 0.29489078 0.31081081    0
## QDA 0.5270270 0.55743243 0.60135135 0.59546464 0.63175676 0.67567568    0
## NB  0.0000000 0.01351351 0.01351351 0.01488338 0.02369493 0.02702703    0
##
## Spec
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## LDA 0.8587571 0.8997175 0.9239351 0.9170507 0.9324890 0.9606742    0
## QDA 0.6440678 0.7090395 0.7211325 0.7185298 0.7299562 0.7683616    0
## NB  0.9887006 0.9957627 1.0000000 0.9977401 1.0000000 1.0000000    0
```

test set performance for Discriminant Analysis

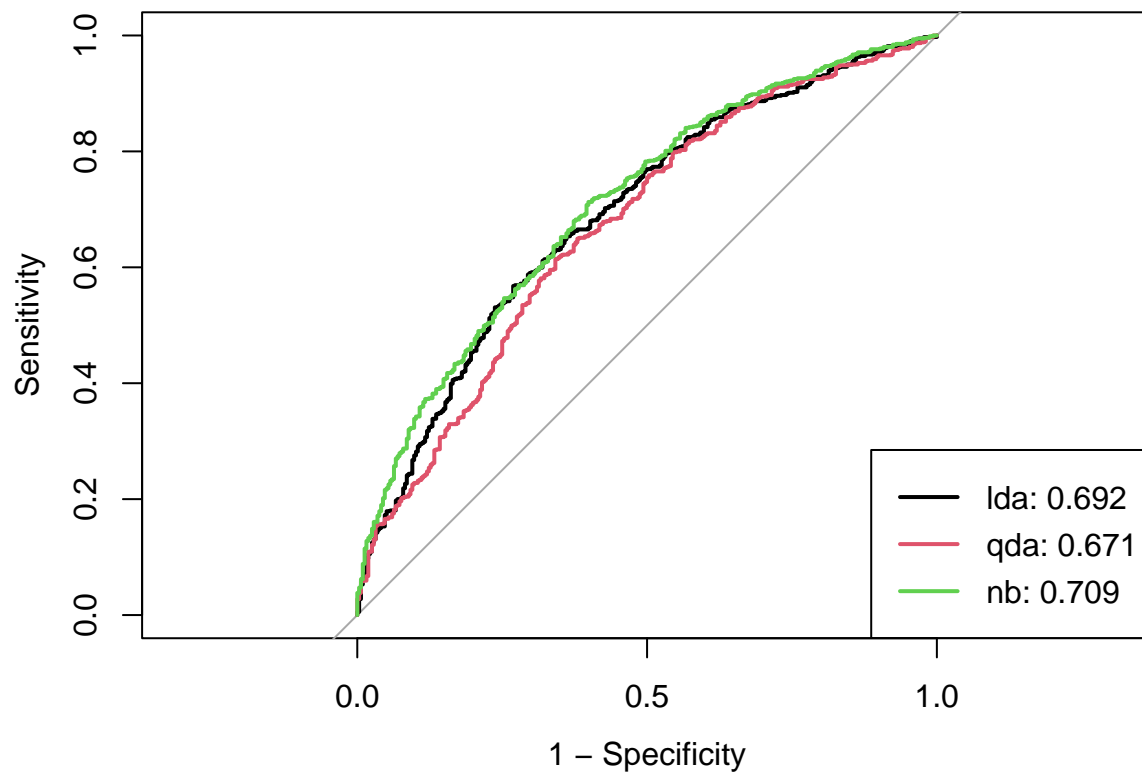
```
lda.pred <- predict(model.lda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
nb.pred <- predict(model.nb, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
qda.pred <- predict(model.qda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]

roc.lda <- roc(covid_dat$recovery_time[-rowTrain], lda.pred)
roc.nb <- roc(covid_dat$recovery_time[-rowTrain], nb.pred)
roc.qda <- roc(covid_dat$recovery_time[-rowTrain], qda.pred)

auc <- c(roc.lda$auc[1], roc.qda$auc[1], roc.nb$auc[1])

plot(roc.lda, legacy.axes = TRUE)
plot(roc.qda, col = 2, add = TRUE)
plot(roc.nb, col = 3, add = TRUE)

modelNames <- c("lda", "qda", "nb")
legend("bottomright", legend = paste0(modelNames, ": ", round(auc, 3)),
      col = 1:3, lwd = 2)
```



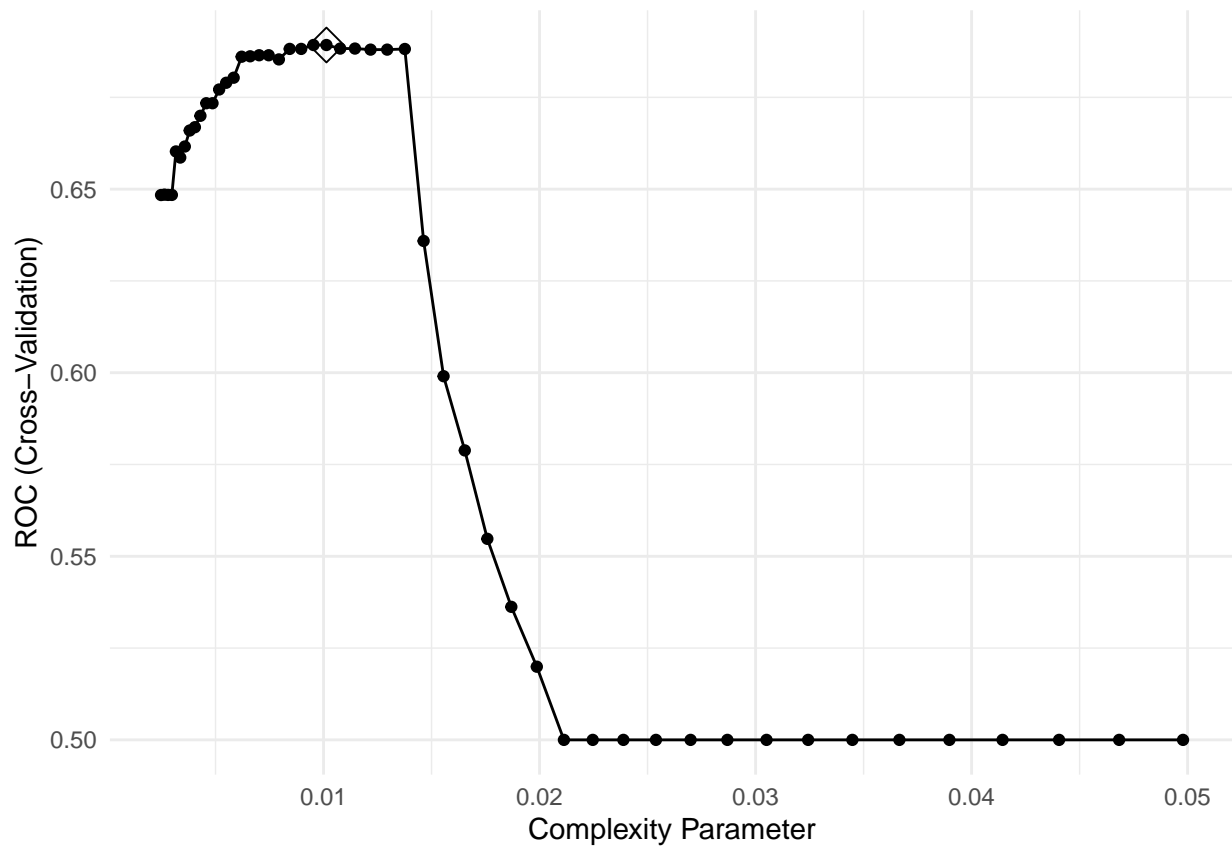
classification tree models

rpart

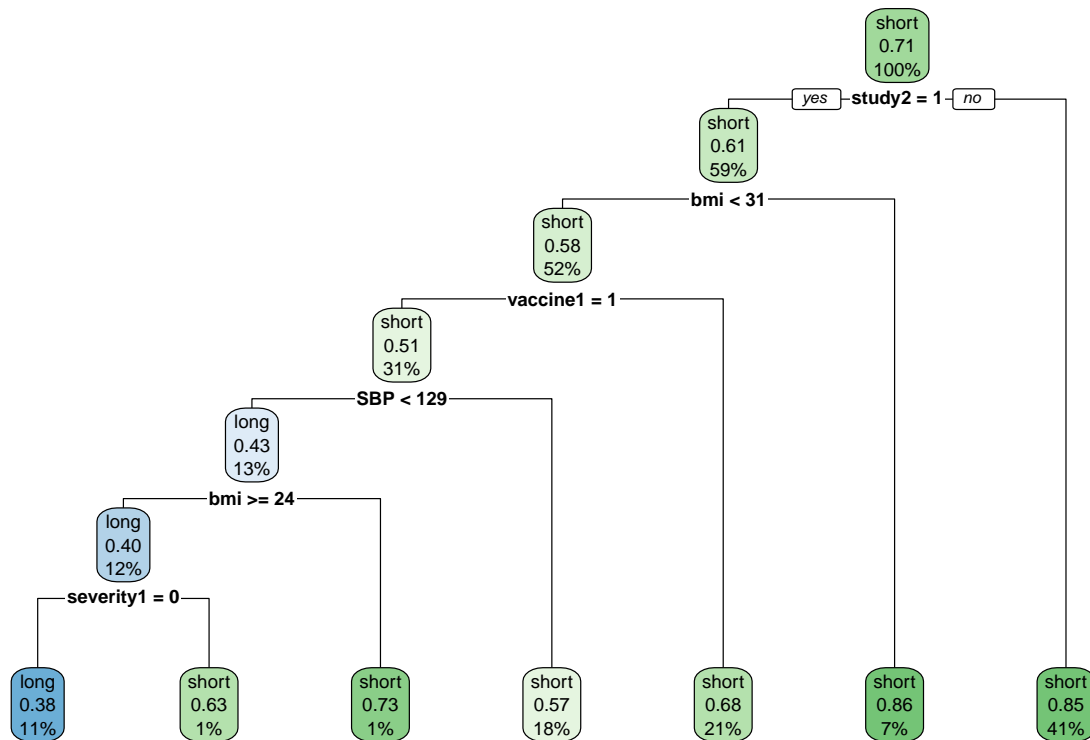
```
num_cores <- detectCores()
cl <- makePSOCKcluster(num_cores)
registerDoParallel(cl)
set.seed(2)

model.rpart = train(recovery_time ~ .,
                    covid_dat,
                    subset = rowTrain,
                    method = "rpart",
                    tuneGrid = data.frame(cp = exp(seq(-6, -3, len = 50))),
                    trControl = ctrl2,
                    metric = "ROC")

ggplot(model.rpart, highlight = TRUE)
```



```
rpart.plot(model.rpart$finalModel)
```



```
stopCluster(c1)
```

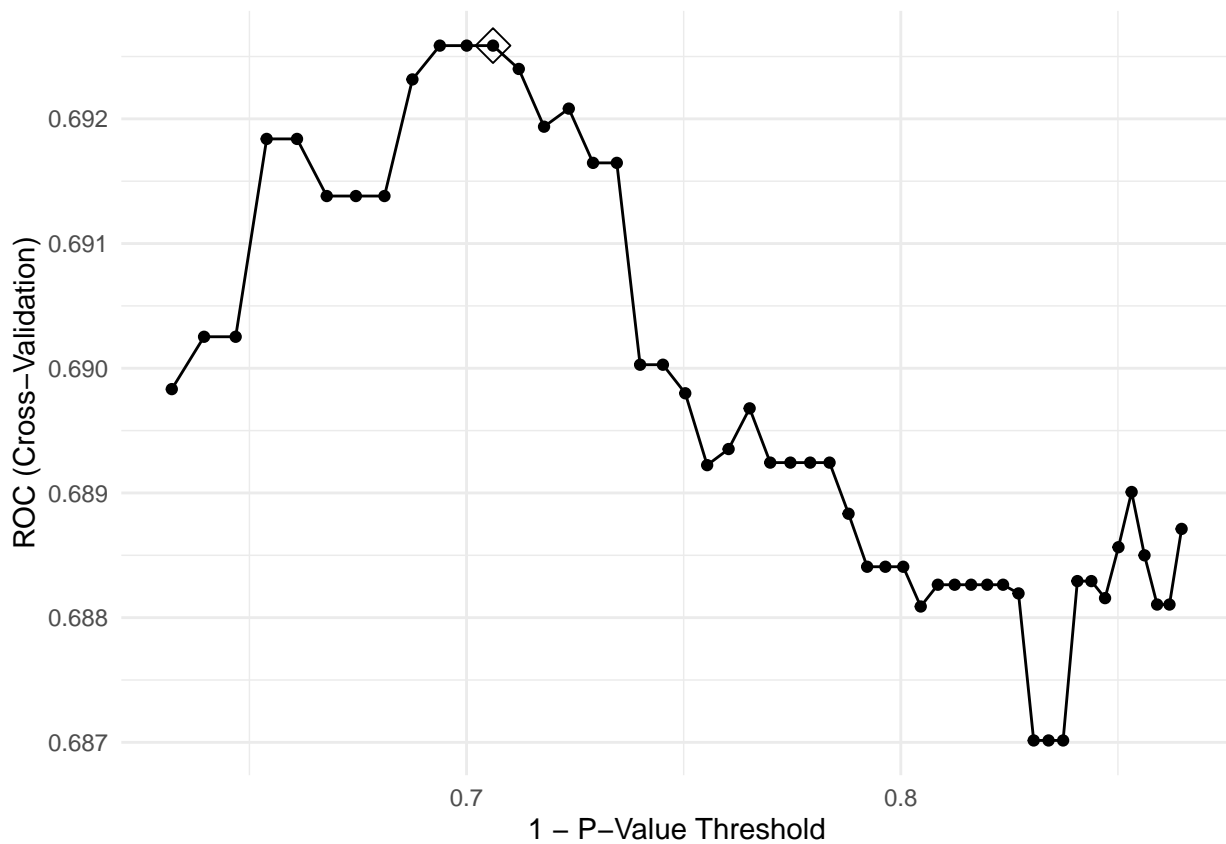
```
registerDoSEQ()
```

ctree


```
num_cores <- detectCores()
cl <- makePSOCKcluster(num_cores)
registerDoParallel(cl)
set.seed(2)

model.ctree = train(recovery_time ~ .,
  covid_dat,
  subset = rowTrain,
  method = "ctree",
  tuneGrid = data.frame(mincriterion = 1 - exp(seq(-2, -1, length = 50))),
  metric = "ROC",
  trControl = ctrl12)

ggplot(model.ctree, highlight = TRUE)
```



```
plot(model.ctree$finalModel)
```

```
stopCluster(cl)
registerDoSEQ()
```

test set performance for classification tree models

```
resamp_tree <- resamples(list(rpart = model.rpart,
                             ctree = model.ctree))
summary(resamp_tree)
```

```
##
## Call:
## summary.resamples(object = resamp_tree)
##
## Models: rpart, ctree
## Number of resamples: 10
##
## ROC
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## rpart 0.6570851 0.6730605 0.6885818 0.6892499 0.6958410 0.7433196    0
## ctree 0.6718201 0.6806484 0.6832281 0.6925873 0.7081043 0.7224385    0
##
## Sens
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## rpart 0.2027027 0.2432432 0.2500000 0.2476490 0.2593947 0.2837838    0
## ctree 0.1621622 0.1790541 0.1959459 0.2288227 0.2837838 0.3287671    0
##
## Spec
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
```

```
## rpart 0.8644068 0.8956151 0.9180791 0.9125627 0.9268631 0.9548023 0
## ctree 0.8474576 0.8997175 0.9152542 0.9103314 0.9255618 0.9604520 0
```

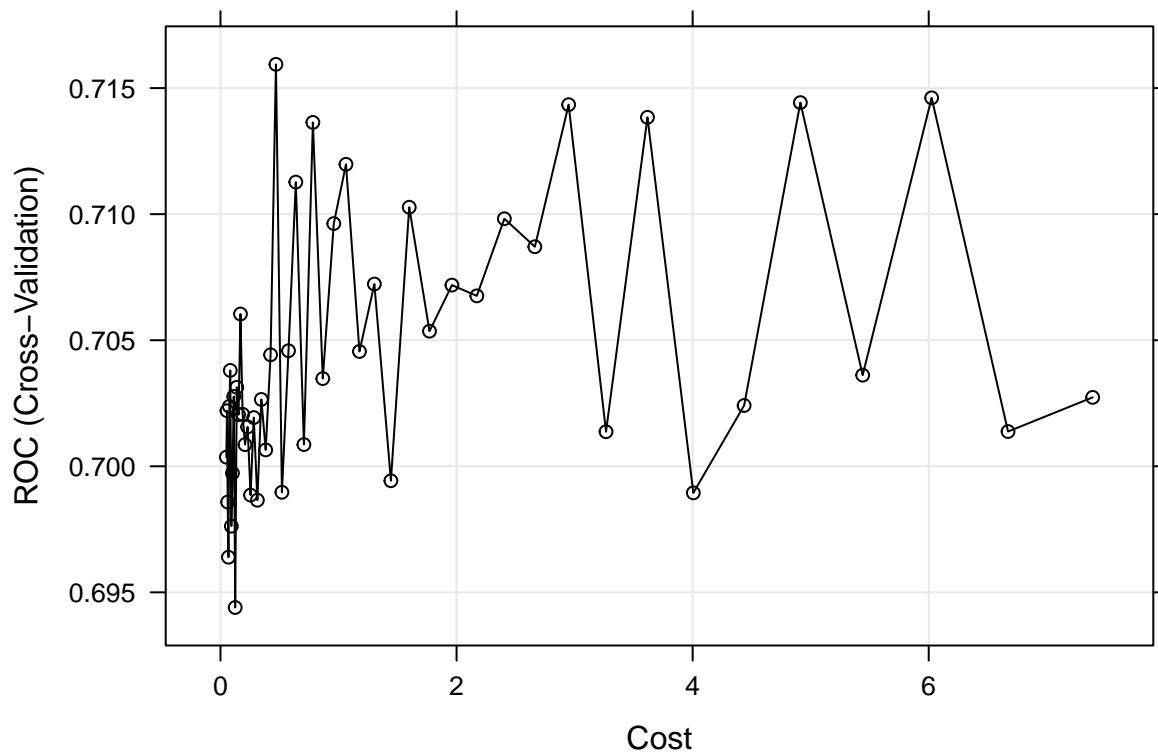
Support Vector Machines

```
num_cores <- detectCores()
cl <- makePSOCKcluster(num_cores)
registerDoParallel(cl)
set.seed(2)
# kernal linear

#model.svm1 <- train(recovery_time ~ .,
#                    #data = covid_dat[rowTrain, ],
#                    #method = "svmLinear",
#                    #tuneGrid = data.frame(C = exp(seq(-2,5,len=5))),
#                    #trControl = ctrl2)

model.svm1 <- train(recovery_time ~ .,
                    data = covid_dat[rowTrain, ],
                    method = "svmLinear2",
                    preProcess = c("center", "scale"),
                    tuneGrid = data.frame(cost = exp(seq(-3,2,len = 50))),
                    trControl = ctrl2)

plot(model.svm1, highlight = TRUE)
```



```
model.svm1$bestTune
```

```
##          cost
```

```

## 23 0.4699627
model.svm1$finalModel

##
## Call:
## svm.default(x = as.matrix(x), y = y, kernel = "linear", cost = param$cost,
##     probability = classProbs)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##     cost:  0.4699627
##
## Number of Support Vectors: 1706
#test error
linear_test_preds = predict(model.svm1, newdata = covid_dat[-rowTrain, ])
confusionMatrix(data = linear_test_preds,
                 reference = covid_dat$recovery_time[-rowTrain])

## Confusion Matrix and Statistics
##
##           Reference
## Prediction long short
##    long    0    0
##    short 316  759
##
##              Accuracy : 0.706
##              95% CI : (0.6778, 0.7331)
##    No Information Rate : 0.706
##    P-Value [Acc > NIR] : 0.5152
##
##              Kappa : 0
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.000
##              Specificity : 1.000
##    Pos Pred Value :    NaN
##    Neg Pred Value : 0.706
##    Prevalence : 0.294
##    Detection Rate : 0.000
##    Detection Prevalence : 0.000
##    Balanced Accuracy : 0.500
##
##    'Positive' Class : long
##
stopCluster(cl)
registerDoSEQ()

num_cores <- detectCores()
cl <- makePSOCKcluster(num_cores)
registerDoParallel(cl)

```

```


svmr.grid <- expand.grid(C = exp(seq(-4,4,len=20)),
                        sigma = exp(seq(-4,0,len=10)))

#radial kernel
set.seed(2)

model.svmr <- train(recovery_time ~ .,
                    data = covid_dat[rowTrain, ],
                    method = "svmRadialSigma",
                    preProcess = c("center", "scale"),
                    tuneGrid = svmr.grid,
                    trControl = ctrl2)

myCol<- rainbow(20)
myPar <- list(superpose.symbol = list(col = myCol),
              superpose.line = list(col = myCol))
ggplot(model.svmr, highlight = TRUE, par.settings = myPar)

```



```

# test error
radial_test_preds = predict(model.svmr, newdata = covid_dat[-rowTrain, ])
confusionMatrix(data = radial_test_preds ,
                 reference = covid_dat$recovery_time[-rowTrain])

```

```

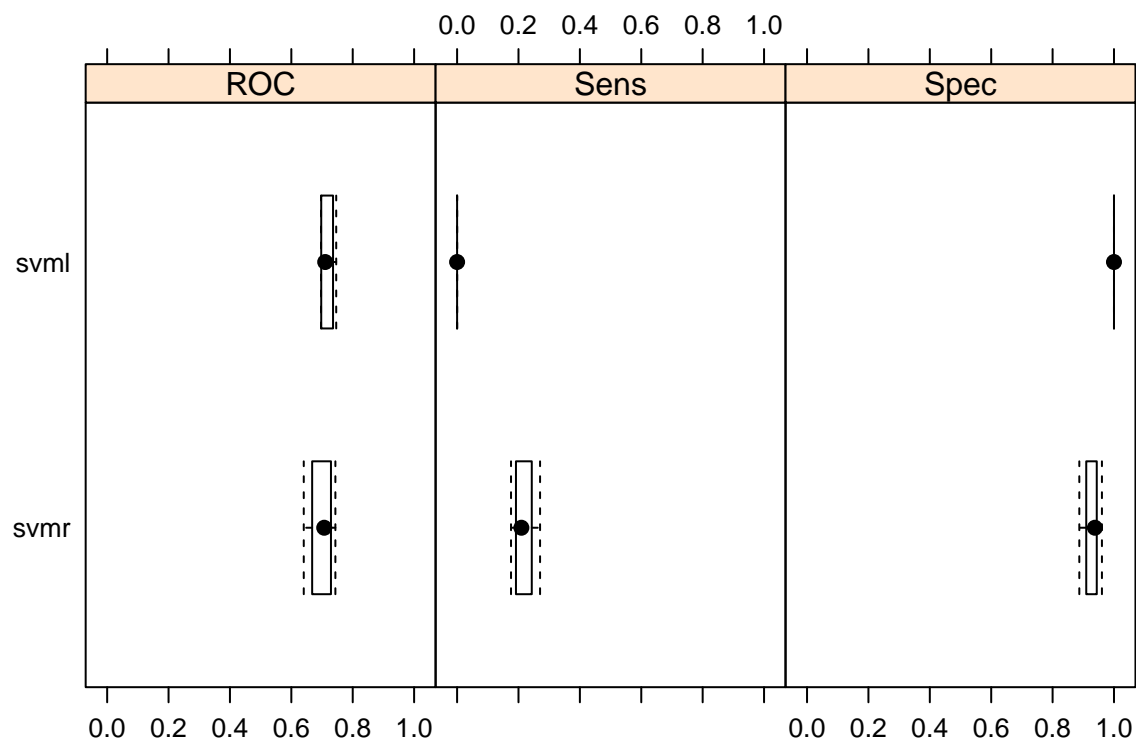
## Confusion Matrix and Statistics
##
##           Reference
## Prediction long short
##      long      77      67
##      short    239     692

```

```
##
##           Accuracy : 0.7153
##           95% CI : (0.6873, 0.7422)
##    No Information Rate : 0.706
##    P-Value [Acc > NIR] : 0.2633
##
##           Kappa : 0.1847
##
##    McNemar's Test P-Value : <2e-16
##
##           Sensitivity : 0.24367
##           Specificity : 0.91173
##           Pos Pred Value : 0.53472
##           Neg Pred Value : 0.74329
##           Prevalence : 0.29395
##           Detection Rate : 0.07163
##    Detection Prevalence : 0.13395
##           Balanced Accuracy : 0.57770
##
##           'Positive' Class : long
##
stopCluster(cl)
registerDoSEQ()
```

test data performance of SVM methods

```
resamp <- resamples(list(svm1 = model.svm1,
                        svmr = model.svmr
                        ))
bwplot(resamp)
```



Model comparison

```
res <- resamples(list(GLM = model.glm, GLMNET = model.glmn, GAM = model.gam, MARS = model.mars, CTREE =
trainROC <- bwplot(res, metric = "ROC")
summary(res)
```

```
##
## Call:
## summary.resamples(object = res)
##
## Models: GLM, GLMNET, GAM, MARS, CTREE, RPART, LDA, QDA, NB, SVML, SVMR
## Number of resamples: 10
##
## ROC
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM      0.7027786 0.7100512 0.7240580 0.7243182 0.7396824 0.7482539    0
## GLMNET   0.7010325 0.7091541 0.7199426 0.7230782 0.7374113 0.7488612    0
## GAM      0.7092686 0.7156054 0.7316766 0.7319728 0.7432725 0.7592621    0
## MARS     0.7131036 0.7220759 0.7309513 0.7309203 0.7385390 0.7542894    0
## CTREE    0.6718201 0.6806484 0.6832281 0.6925873 0.7081043 0.7224385    0
## RPART    0.6570851 0.6730605 0.6885818 0.6892499 0.6958410 0.7433196    0
## LDA      0.6982994 0.7092113 0.7204289 0.7236334 0.7385363 0.7509869    0
## QDA      0.6825470 0.6870515 0.7021897 0.7067165 0.7212097 0.7436988    0
## NB       0.6882730 0.7120171 0.7177622 0.7186405 0.7223240 0.7519739    0
## SVML     0.6968239 0.6982173 0.7103583 0.7159424 0.7324267 0.7462041    0
## SVMR     0.6404795 0.6755612 0.7070927 0.6985938 0.7242925 0.7435469    0
##
## Sens
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
```

```

## GLM      0.2162162 0.27027027 0.27897075 0.28146983 0.30743243 0.33783784      0
## GLMNET   0.2027027 0.24662162 0.28378378 0.26932618 0.29697334 0.32432432      0
## GAM      0.2027027 0.26277305 0.31756757 0.30170307 0.33783784 0.36486486      0
## MARS     0.2297297 0.28378378 0.29248427 0.30444280 0.33445946 0.37837838      0
## CTREE    0.1621622 0.17905405 0.19594595 0.22882266 0.28378378 0.32876712      0
## RPART    0.2027027 0.24324324 0.25000000 0.24764902 0.25939467 0.28378378      0
## LDA      0.1891892 0.26013514 0.27702703 0.26525361 0.29489078 0.31081081      0
## QDA      0.5270270 0.55743243 0.60135135 0.59546464 0.63175676 0.67567568      0
## NB       0.0000000 0.01351351 0.01351351 0.01488338 0.02369493 0.02702703      0
## SVM      0.0000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000      0
## SVMR     0.1756757 0.19451129 0.20945946 0.21917808 0.24324324 0.27027027      0
##
## Spec
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## GLM      0.8531073 0.9039548 0.9154605 0.9119755 0.9324890 0.9438202      0
## GLMNET   0.8644068 0.9053672 0.9180791 0.9187425 0.9382022 0.9548023      0
## GAM      0.8644068 0.8884181 0.9124294 0.9018314 0.9157303 0.9269663      0
## MARS     0.8644068 0.8912429 0.9098584 0.9057862 0.9196106 0.9325843      0
## CTREE    0.8474576 0.8997175 0.9152542 0.9103314 0.9255618 0.9604520      0
## RPART    0.8644068 0.8956151 0.9180791 0.9125627 0.9268631 0.9548023      0
## LDA      0.8587571 0.8997175 0.9239351 0.9170507 0.9324890 0.9606742      0
## QDA      0.6440678 0.7090395 0.7211325 0.7185298 0.7299562 0.7683616      0
## NB       0.9887006 0.9957627 1.0000000 0.9977401 1.0000000 1.0000000      0
## SVM      1.0000000 1.0000000 1.0000000 1.0000000 1.0000000 1.0000000      0
## SVMR     0.8870056 0.9124294 0.9378531 0.9288929 0.9423284 0.9606742      0

# Prediction on test set
glm.pred <- predict(model.glm, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
glmn.pred <- predict(model.glmn, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
gam.pred <- predict(model.gam, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
mars.pred <- predict(model.mars, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
ctree.pred <- predict(model.ctree, newdata = covid_dat[-rowTrain,], type = "prob")[,2]
rpart.pred <- predict(model.rpart, newdata = covid_dat[-rowTrain,], type = "prob")[,2]
lda.pred <- predict(model.lda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
qda.pred <- predict(model.qda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
nb.pred <- predict(model.nb, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
svml.pred <- predict(model.svml, newdata = covid_dat[-rowTrain,], type = "prob")[,2]
svmr.pred <- predict(model.svmr, newdata = covid_dat[-rowTrain,], type = "prob")[,2]

# Test error rate for all models

test_glm_error = mean(glm.pred != covid_dat$recovery_time[-rowTrain])
test_glmn_error = mean(glmn.pred != covid_dat$recovery_time[-rowTrain])
test_gam_error = mean(gam.pred != covid_dat$recovery_time[-rowTrain])
test_mars_error = mean(mars.pred != covid_dat$recovery_time[-rowTrain])
test_ctree_error = mean(ctree.pred != covid_dat$recovery_time[-rowTrain])
test_rpart_error = mean(rpart.pred != covid_dat$recovery_time[-rowTrain])
test_lda_error = mean(lda.pred != covid_dat$recovery_time[-rowTrain])
test_qda_error = mean(qda.pred != covid_dat$recovery_time[-rowTrain])
test_nb_error = mean(nb.pred != covid_dat$recovery_time[-rowTrain])
test_svml_error = mean(svml.pred != covid_dat$recovery_time[-rowTrain])
test_svmr_error = mean(svmr.pred != covid_dat$recovery_time[-rowTrain])

# Generalized Linear Model (GLM)

```

```

test_glm_error

## [1] 1
# Penalized logistic regression (GLMN)
test_glmn_error

## [1] 1
# Generalized Additive Model (GAM)
test_gam_error

## [1] 1
# Multivariate Adaptive Regression Splines (MARS)
test_mars_error

## [1] 1
# Conditional Inference Tree (CART)
test_ctree_error

## [1] 1
# Recursive Partitioning and Regression Trees (RPART)
test_rpart_error

## [1] 1
# Linear Discriminant Analysis (LDA)
test_lda_error

## [1] 1
# Quadratic Discriminant Analysis (QDA)
test_qda_error

## [1] 1
# Naive Bayes (NB)
test_nb_error

## [1] 1
# Support Vector Machine Regression (SVMR)
pred_svmr_test = predict(model.svmr, newdata=covid_dat[-rowTrain,], type="raw")
test_svmr_error = mean(pred_svmr_test != covid_dat$recovery_time[-rowTrain])
test_svmr_error

## [1] 0.2846512
# Support Vector Machine Classification (SVML)
pred_svml_test = predict(model.svml, newdata=covid_dat[-rowTrain,], type="raw")
test_svml_error = mean(pred_svml_test != covid_dat$recovery_time[-rowTrain])
test_svml_error

## [1] 0.2939535
#create the table
model_errors <- data.frame(Model = c("GLM", "GLMN", "GAM", "MARS", "CART", "RPART", "LDA", "QDA", "NB",
Error_Rate = c(test_glm_error, test_glmn_error, test_gam_error, test_mars_error,
test_ctree_error, test_rpart_error, test_lda_error, test_qda_error, test_nb_error, test_svmr_error, tes

```



```
knitr::kable(model_errors, align = "c")
```

Model	Error_Rate
GLM	1.0000000
GLMN	1.0000000
GAM	1.0000000
MARS	1.0000000
CART	1.0000000
RPART	1.0000000
LDA	1.0000000
QDA	1.0000000
NB	1.0000000
SVMR	0.2846512
SVML	0.2939535

```
#roc
roc.glm <- roc(covid_dat$recovery_time[-rowTrain], glm.pred)
roc.glmn <- roc(covid_dat$recovery_time[-rowTrain], glmn.pred)
roc.gam <- roc(covid_dat$recovery_time[-rowTrain], gam.pred)
roc.mars <- roc(covid_dat$recovery_time[-rowTrain], mars.pred)

roc.ctree <- roc(covid_dat$recovery_time[-rowTrain], ctree.pred)
roc.rpart <- roc(covid_dat$recovery_time[-rowTrain], rpart.pred)

roc.lda <- roc(covid_dat$recovery_time[-rowTrain], lda.pred)
roc.qda <- roc(covid_dat$recovery_time[-rowTrain], qda.pred)
roc.nb <- roc(covid_dat$recovery_time[-rowTrain], nb.pred)

roc.svml <- roc(covid_dat$recovery_time[-rowTrain], svml.pred)
roc.svmr <- roc(covid_dat$recovery_time[-rowTrain], svmr.pred)

auc <- c(roc.glm$auc[1], roc.glmn$auc[1],
         roc.gam$auc[1], roc.mars$auc[1],
         roc.lda$auc[1], roc.qda$auc[1], roc.nb$auc[1],
         roc.ctree$auc[1], roc.rpart$auc[1],
         roc.svml$auc[1], roc.svmr$auc[1]
        )

plot(roc.glm, legacy.axes = TRUE)
plot(roc.glmn, col = 2, add = TRUE)


plot(roc.gam, col = 3, add = TRUE)
plot(roc.mars, col = 4, add = TRUE)

plot(roc.lda, col = 5, add = TRUE)
plot(roc.qda, col = 6, add = TRUE)
plot(roc.nb, col = 7, add = TRUE)

plot(roc.ctree, col = 8, add = TRUE)
plot(roc.rpart, col = 9, add = TRUE)
```

```
plot(roc.svm1, col = 10, add = TRUE)
plot(roc.svmr, col = 11, add = TRUE)
```

```
modelNames <- c("glm","glmnet","gam","mars","lda","qda","nb", "ctree", "rpart","svm (linear kernel)", "svm (radial kernel)")
legend("bottomright", legend = paste0(modelNames, ":", round(auc,3)),
      col = 1:11, lwd = 2)
```



```
# interpretation for final mars model
model.mars$finalModel
```

```
## GLM (family binomial, link logit):
## nulldev df      dev df   devratio    AIC iters converged
## 3043.84 2511    2649.02 2501     0.13    2671     5           1
##
## Earth selected 11 of 17 terms, and 9 of 18 predictors (nprune=11)
## Termination condition: RSq changed by less than 0.001 at 17 terms
## Importance: study2, bmi, vaccine1, SBP, severity1, smoking1, gender1, ...
## Number of terms at each degree of interaction: 1 10 (additive model)
## Earth GCV 0.1815234    RSS 448.3946    GRSq 0.1264802    RSq 0.1403399
summary(model.mars$finalModel)
```


```
## Call: earth(x=matrix[2512,18], y=factor.object, keepxy=TRUE,
##           glm=list(family=function(object, maxit=100), degree=1, nprune=11)
##
## GLM coefficients
##               short
## (Intercept) -0.32524568
## gender1     -0.32207625
```

```
## smoking1      0.43021337
## smoking2      0.55022116
## vaccine1     -0.73109733
## severity1      0.80307433
## study2       -1.35310824
## h(bmi-23.1)   0.41456148
## h(28.6-bmi)  0.51047027
## h(135-SBP)   -0.03262848
## h(LDL-145)   -0.05342548
##
## GLM (family binomial, link logit):
## nulldev  df      dev  df  devratio    AIC iters converged
## 3043.84 2511  2649.02 2501      0.13   2671      5           1
##
## Earth selected 11 of 17 terms, and 9 of 18 predictors (nprune=11)
## Termination condition: RSq changed by less than 0.001 at 17 terms
## Importance: study2, bmi, vaccine1, SBP, severity1, smoking1, gender1, ...
## Number of terms at each degree of interaction: 1 10 (additive model)
## Earth GCV 0.1815234    RSS 448.3946    GRSq 0.1264802    RSq 0.1403399
```

```
model.mars$bestTune
```

```
##      nprune degree
## 10      11      1
```

```
plot(model.mars)
```



```
coef(model.mars$finalModel)
```

```
## (Intercept)      study2 h(28.6-bmi)    vaccine1 h(135-SBP)    severity1
## -0.32524568 -1.35310824  0.51047027 -0.73109733 -0.03262848  0.80307433
```

```
##      smoking1      gender1      smoking2 h(LDL-145) h(bmi-23.1)
## 0.43021337 -0.32207625 0.55022116 -0.05342548 0.41456148
vip(model.mars$finalModel) + ggtitle("Gradient Boosting")
```

