

P8106_yiminchen_secondaryanalysis

Yimin Chen (yc4195), Yang Yi (yy3307), Qingyue Zhuo (qz2493)

Contents

Import and data manipulation	1
Data visualization	3
Model training	3
Logistic regression and its cousins	3
GLM	3
Penalized logistic regression	3
GAM	4
MARS	5
test data performance for Logistic regression and its cousins	8
Discriminant Analysis	9
LDA	9
QDA	9
Naive Bayes (NB)	9
test set performance for Discriminant Analysis	11
classification tree models	12
rpart	12
ctree	14
test set performance for classification tree models	15
Support Vector Machines	16

Import and data manipulation

```
# Load recovery.RData environment
load("./recovery.Rdata")

dat %>% na.omit()

# dat1 draw a random sample of 2000 participants Uni:3307
set.seed(3307)

dat1 = dat[sample(1:10000, 2000),]

dat1 =
  dat1[, -1] %>%
  mutate(
    recovery_time = as.factor(
      case_when(recovery_time <= 30 ~ "long", recovery_time > 30 ~ "short")
    ),
    gender = as.factor(gender),
```

```

    race = as.factor(race),
    smoking = as.factor(smoking),
    hypertension = as.factor(hypertension),
    diabetes = as.factor(diabetes),
    vaccine = as.factor(vaccine),
    severity = as.factor(severity),
    study = as.factor(
      case_when(study == "A" ~ 1, study == "B" ~ 2, study == "C" ~ 3)
    )
  )
)

# dat2 draw a random sample of 2000 participants Uni:2493
set.seed(2493)

dat2 = dat[sample(1:10000, 2000),]

dat2 =
  dat2[, -1] %>%
  mutate(
    recovery_time = as.factor(
      case_when(recovery_time <= 30 ~ "long", recovery_time > 30 ~ "short")
    ),
    gender = as.factor(gender),
    race = as.factor(race),
    smoking = as.factor(smoking),
    hypertension = as.factor(hypertension),
    diabetes = as.factor(diabetes),
    vaccine = as.factor(vaccine),
    severity = as.factor(severity),
    study = as.factor(
      case_when(study == "A" ~ 1, study == "B" ~ 2, study == "C" ~ 3)
    )
  )

# Merged dataset with unique observation
covid_dat = rbind(dat1, dat2) %>%
  unique()

covid_dat2 = model.matrix(recovery_time ~ ., covid_dat)[, -1] #ignore intercept

# Partition dataset into two parts: training data (70%) and test data (30%)
rowTrain = createDataPartition(y = covid_dat$recovery_time, p = 0.7, list = FALSE)

trainData = covid_dat[rowTrain, ]
testData = covid_dat[-rowTrain, ]

# matrix of predictors
x1 = covid_dat2[rowTrain,]
# vector of response
y1 = covid_dat$recovery_time[rowTrain]
# matrix of predictors
x2 = covid_dat2[-rowTrain,]

```

```
# vector of response
y2 = covid_dat$recovery_time[-rowTrain]

ctrl1 = trainControl(method = "repeatedcv", number = 10, repeats = 5)
ctrl2 = trainControl(method = "cv",
                      classProbs = TRUE,
                      summaryFunction = twoClassSummary)
```

Data visualization

Model training

classification

- glm + penalized logistic regression L8
- GAM L8
- MARS L8
- QDA L9
- LDA L9
- Naive Bayes L9
- classification tree: L11
- random forest L12
- boosting L12
- support vector machines L13

Logistic regression and its cousins

GLM

```
set.seed(2)
model.glm <- train(x = covid_dat2[rowTrain,],
                   y = covid_dat$recovery_time[rowTrain],
                   method = "glm",
                   metric = "ROC",
                   trControl = ctrl2)
```

Penalized logistic regression

Penalized logistic regression can be fitted using `glmnet`. We use the `train` function to select the optimal tuning parameters.

```
glmnetGrid <- expand.grid(.alpha = seq(0, 1, length = 21),
                        .lambda = exp(seq(-8, -1, length = 50)))

set.seed(2)
model.glmnet <- train(x = covid_dat2[rowTrain,],
                     y = covid_dat$recovery_time[rowTrain],
                     method = "glmnet",
                     tuneGrid = glmnetGrid,
                     metric = "ROC",
                     trControl = ctrl2)


model.glmnet$bestTune
```

```
##      alpha      lambda
```

```
## 1001      1 0.0003354626
```

```
myCol<- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
              superpose.line = list(col = myCol))

plot(model.glmn, par.settings = myPar, xTrans = function(x) log(x))
```



GAM


```
set.seed(2)
model.gam <- train(x = covid_dat2[rowTrain,],
                   y = covid_dat$recovery_time[rowTrain],
                   method = "gam",
                   metric = "ROC",
                   trControl = ctrl2)
```

```
model.gam$finalModel
```

```
##
## Family: binomial
## Link function: logit
##
## Formula:
## .outcome ~ gender1 + race3 + race4 + smoking1 + smoking2 + hypertension1 +
##           diabetes1 + vaccine1 + severity1 + study2 + study3 + s(age) +
##           s(SBP) + s(LDL) + s(bmi) + s(height) + s(weight)
##
```

```
## Estimated degrees of freedom:
## 1.00 1.00 1.55 2.79 1.00 2.29 total = 21.63
##
## UBRE score: 0.06815373
```


```
plot(model.gam$finalModel, select = 3)
```



MARS

```
set.seed(2)
model.mars <- train(x = covid_dat2[rowTrain,],
                    y = covid_dat$recovery_time[rowTrain],
                    method = "earth",
                    tuneGrid = expand.grid(degree = 1:4,
                                           nprune = 2:20),
                    metric = "ROC",
                    trControl = ctrl12)
```


```
plot(model.mars)
```



```
coef(model.mars$finalModel)
```

```
## (Intercept)      study2 h(28.6-bmi)    vaccine1 h(135-SBP)    severity1
## -0.32524568 -1.35310824  0.51047027 -0.73109733 -0.03262848  0.80307433
##      smoking1      gender1      smoking2 h(LDL-145) h(bmi-23.1)
##  0.43021337 -0.32207625  0.55022116 -0.05342548  0.41456148
```

```
vip(model.mars$finalModel)
```




```
res <- resamples(list(GLM = model.glm,
                      GLMNET = model.glmn,
                      GAM = model.gam,
                      MARS = model.mars))
summary(res)
```

```
##
## Call:
## summary.resamples(object = res)
##
## Models: GLM, GLMNET, GAM, MARS
## Number of resamples: 10
##
## ROC
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM      0.7027786 0.7100512 0.7240580 0.7243182 0.7396824 0.7482539    0
## GLMNET    0.7010325 0.7091541 0.7199426 0.7230782 0.7374113 0.7488612    0
## GAM      0.7092686 0.7156054 0.7316766 0.7319728 0.7432725 0.7592621    0
## MARS      0.7131036 0.7220759 0.7309513 0.7309203 0.7385390 0.7542894    0
##
## Sens
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM      0.2162162 0.2702703 0.2789708 0.2814698 0.3074324 0.3378378    0
## GLMNET    0.2027027 0.2466216 0.2837838 0.2693262 0.2969733 0.3243243    0
## GAM      0.2027027 0.2627730 0.3175676 0.3017031 0.3378378 0.3648649    0
## MARS      0.2297297 0.2837838 0.2924843 0.3044428 0.3344595 0.3783784    0
##
```

```
## Spec
##           Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## GLM      0.8531073 0.9039548 0.9154605 0.9119755 0.9324890 0.9438202    0
## GLMNET 0.8644068 0.9053672 0.9180791 0.9187425 0.9382022 0.9548023    0
## GAM      0.8644068 0.8884181 0.9124294 0.9018314 0.9157303 0.9269663    0
## MARS     0.8644068 0.8912429 0.9098584 0.9057862 0.9196106 0.9325843    0
```

```
bwplot(res, metric = "ROC")
```



test data performance for Logistic regression and its cousins


```
glm.pred <- predict(model.glm, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
glmnet.pred <- predict(model.glmnet, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
gam.pred <- predict(model.gam, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
mars.pred <- predict(model.mars, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]

roc.glm <- roc(covid_dat$recovery_time[-rowTrain], glm.pred)
roc.glmnet <- roc(covid_dat$recovery_time[-rowTrain], glmnet.pred)
roc.gam <- roc(covid_dat$recovery_time[-rowTrain], gam.pred)
roc.mars <- roc(covid_dat$recovery_time[-rowTrain], mars.pred)

auc <- c(roc.glm$auc[1], roc.glmnet$auc[1],
         roc.gam$auc[1], roc.mars$auc[1])

modelNames <- c("glm", "glmnet", "gam", "mars")

ggroc(list(roc.glm, roc.glmnet, roc.gam, roc.mars), legacy.axes = TRUE) +
  scale_color_discrete(labels = paste0(modelNames, " (", round(auc, 3), ")"),
                       name = "Models (AUC)") +
  geom_abline(intercept = 0, slope = 1, color = "grey")
```

Discriminant Analysis

LDA

```
set.seed(2)

model.lda <- train(x = covid_dat2[rowTrain,],
  y = covid_dat$recovery_time[rowTrain],
  method = "lda",
  metric = "ROC",
  trControl = ctrl2)
```

QDA

```
set.seed(2)

model.qda <- train(x = covid_dat2[rowTrain,],
  y = covid_dat$recovery_time[rowTrain],
  method = "qda",
  metric = "ROC",
  trControl = ctrl2)
```

Naive Bayes (NB)


There is one practical issue with the NB classifier when nonparametric estimators are used. When a new data point includes a feature value that never occurs for some response class, the posterior probability can become zero. To avoid this, we increase the count of the value with a zero occurrence to a small value, so

that the overall probability doesn't become zero. In practice, a value of one or two is a common choice. This correction is called "Laplace Correction," and is implemented via the parameter `fL`. The parameter `adjust` adjusts the bandwidths of the kernel density estimates, and a larger value means a more flexible estimate.

```
nbGrid <- expand.grid(usekernel = c(FALSE, TRUE),
                     fL = 1,
                     adjust = seq(.2, 3, by = .2))

set.seed(2)
model.nb <- train(x = covid_dat2[rowTrain,],
                  y = covid_dat$recovery_time[rowTrain],
                  method = "nb",
                  tuneGrid = nbGrid,
                  metric = "ROC",
                  trControl = ctrl12)

plot(model.nb)
```



```
res <- resamples(list(LDA = model.lda, QDA = model.qda, NB = model.nb))
summary(res)
```

```
##
## Call:
## summary.resamples(object = res)
##
## Models: LDA, QDA, NB
## Number of resamples: 10
##
## ROC
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
```

```
## LDA 0.6982994 0.7092113 0.7204289 0.7236334 0.7385363 0.7509869 0
## QDA 0.6825470 0.6870515 0.7021897 0.7067165 0.7212097 0.7436988 0
## NB 0.6882730 0.7120171 0.7177622 0.7186405 0.7223240 0.7519739 0
##
## Sens
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## LDA 0.1891892 0.26013514 0.27702703 0.26525361 0.29489078 0.31081081 0
## QDA 0.5270270 0.55743243 0.60135135 0.59546464 0.63175676 0.67567568 0
## NB 0.0000000 0.01351351 0.01351351 0.01488338 0.02369493 0.02702703 0
##
## Spec
##      Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## LDA 0.8587571 0.8997175 0.9239351 0.9170507 0.9324890 0.9606742 0
## QDA 0.6440678 0.7090395 0.7211325 0.7185298 0.7299562 0.7683616 0
## NB 0.9887006 0.9957627 1.0000000 0.9977401 1.0000000 1.0000000 0
```

test set performance for Discriminant Analysis


```
lda.pred <- predict(model.lda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
nb.pred <- predict(model.nb, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
qda.pred <- predict(model.qda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]

roc.lda <- roc(covid_dat$recovery_time[-rowTrain], lda.pred)
roc.nb <- roc(covid_dat$recovery_time[-rowTrain], nb.pred)
roc.qda <- roc(covid_dat$recovery_time[-rowTrain], qda.pred)

auc <- c(roc.lda$auc[1], roc.qda$auc[1], roc.nb$auc[1])

plot(roc.lda, legacy.axes = TRUE)
plot(roc.qda, col = 2, add = TRUE)
plot(roc.nb, col = 3, add = TRUE)

modelNames <- c("lda", "qda", "nb")
legend("bottomright", legend = paste0(modelNames, ": ", round(auc, 3)),
      col = 1:3, lwd = 2)
```




classification tree models

rpart

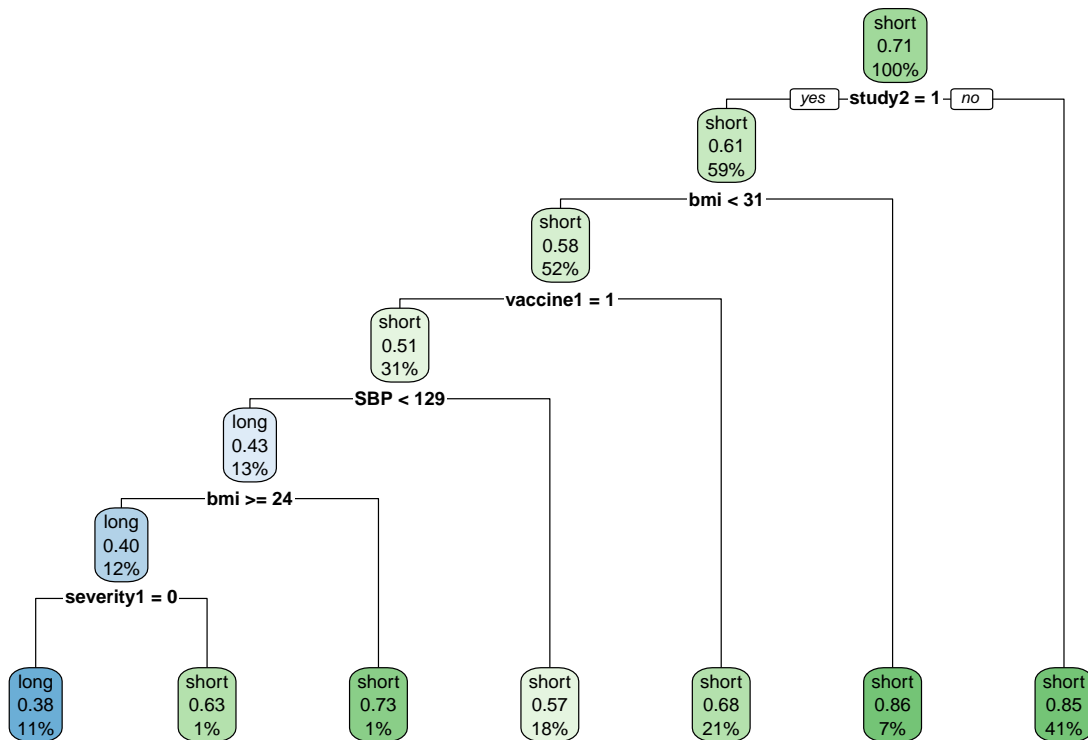
```
set.seed(2)

model.rpart = train(recovery_time ~ .,
  covid_dat,
  subset = rowTrain,
  method = "rpart",
  tuneGrid = data.frame(cp = exp(seq(-6, -3, len = 50))),
  trControl = ctrl2,
  metric = "ROC")

ggplot(model.rpart, highlight = TRUE)
```



```
rpart.plot(model.rpart$finalModel)
```




ctree

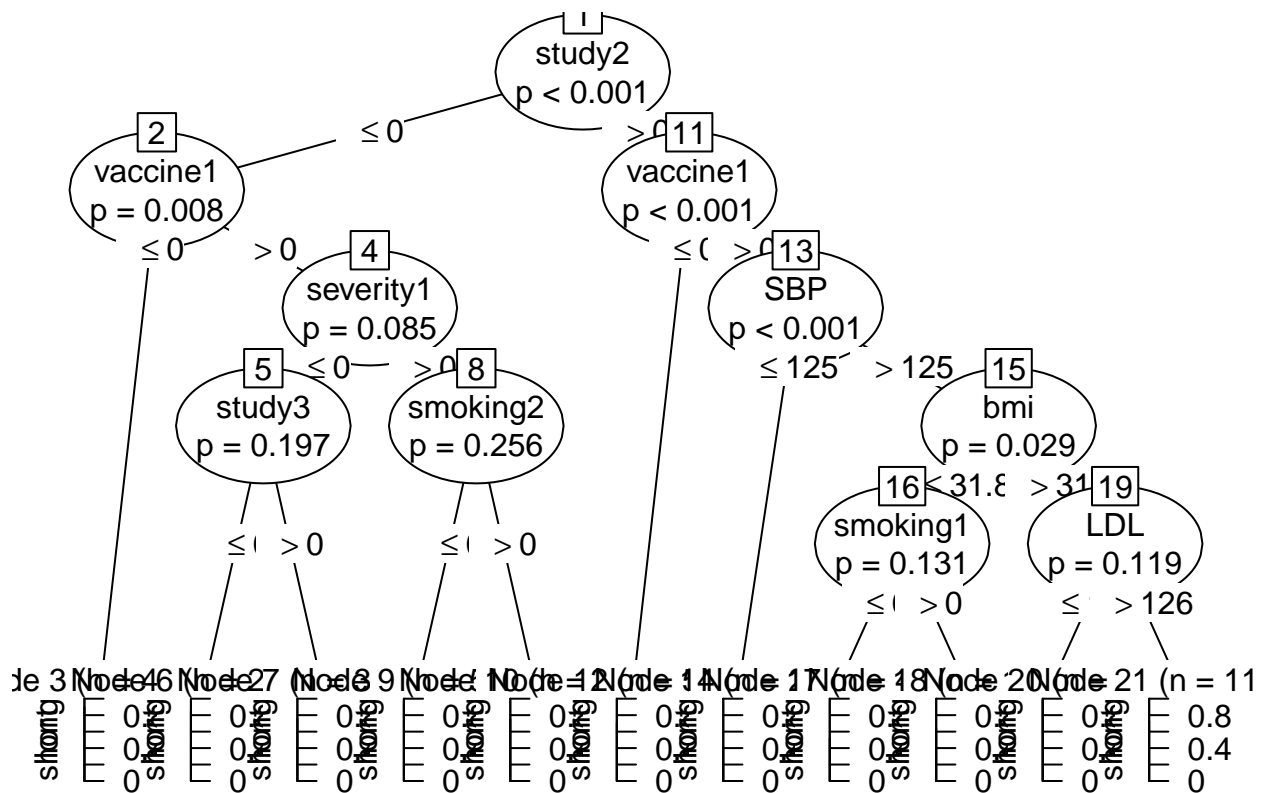
```
set.seed(2)
```

```
model.ctree = train(recovery_time ~ .,  
                    covid_dat,  
                    subset = rowTrain,  
                    method = "ctree",  
                    tuneGrid = data.frame(mincriterion = 1 - exp(seq(-2, -1, length = 50))),  
                    metric = "ROC",  
                    trControl = ctrl2)
```

```
ggplot(model.ctree, highlight = TRUE)
```



```
plot(model.ctree$finalModel)
```



test set performance for classification tree models

```
resamp_tree <- resamples(list(rpart = model.rpart,
                             ctree = model.ctree))
summary(resamp_tree)
```

```
##
## Call:
## summary.resamples(object = resamp_tree)
##
## Models: rpart, ctree
## Number of resamples: 10
##
## ROC
##           Min.    1st Qu.    Median      Mean   3rd Qu.     Max. NA's
## rpart 0.6570851 0.6730605 0.6885818 0.6892499 0.6958410 0.7433196    0
## ctree 0.6718201 0.6806484 0.6832281 0.6925873 0.7081043 0.7224385    0
##
## Sens
##           Min.    1st Qu.    Median      Mean   3rd Qu.     Max. NA's
## rpart 0.2027027 0.2432432 0.2500000 0.2476490 0.2593947 0.2837838    0
## ctree 0.1621622 0.1790541 0.1959459 0.2288227 0.2837838 0.3287671    0
##
## Spec
##           Min.    1st Qu.    Median      Mean   3rd Qu.     Max. NA's
## rpart 0.8644068 0.8956151 0.9180791 0.9125627 0.9268631 0.9548023    0
## ctree 0.8474576 0.8997175 0.9152542 0.9103314 0.9255618 0.9604520    0
```


Support Vector Machines

```
set.seed(2)
# kernal linear

#model.svm1 <- train(recovery_time ~ .,
#                    #data = covid_dat[rowTrain, ],
#                    #method = "svmLinear",
#                    #tuneGrid = data.frame(C = exp(seq(-2,5,len=5))),
#                    #trControl = ctrl2)

model.svm1 <- train(recovery_time ~ .,
                    data = covid_dat[rowTrain, ],
                    method = "svmLinear2",
                    preProcess = c("center", "scale"),
                    tuneGrid = data.frame(cost = exp(seq(-3,2,len = 50))),
                    trControl = ctrl2)

plot(model.svm1, highlight = TRUE)
```



```
model.svm1$bestTune
```

```
##          cost
## 23 0.4699627
```

```
model.svm1$finalModel
```

```
##
## Call:
## svm.default(x = as.matrix(x), y = y, kernel = "linear", cost = param$cost,
```



```
##      probability = classProbs)
##
##
## Parameters:
##      SVM-Type:  C-classification
##      SVM-Kernel:  linear
##      cost:  0.4699627
##
## Number of Support Vectors:  1706
#test error
linear_test_preds = predict(model.svml, newdata = covid_dat[-rowTrain, ])
confusionMatrix(data = linear_test_preds,
                 reference = covid_dat$recovery_time[-rowTrain])
```


```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction long short
##      long      0      0
##      short 316    759
##
##              Accuracy : 0.706
##              95% CI : (0.6778, 0.7331)
##      No Information Rate : 0.706
##      P-Value [Acc > NIR] : 0.5152
##
##              Kappa : 0
##
## Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.000
##              Specificity : 1.000
##      Pos Pred Value :      NaN
##      Neg Pred Value : 0.706
##      Prevalence : 0.294
##      Detection Rate : 0.000
##      Detection Prevalence : 0.000
##      Balanced Accuracy : 0.500
##
##      'Positive' Class : long
##
```

```
svmr.grid <- expand.grid(C = exp(seq(-4,4,len=20)),
                       sigma = exp(seq(-4,0,len=10)))
#radial kernel
set.seed(2)

model.svmr <- train(recovery_time ~ .,
                   data = covid_dat[rowTrain, ],
                   method = "svmRadialSigma",
                   preProcess = c("center", "scale"),
                   tuneGrid = svmr.grid,
                   trControl = ctrl2)
```

```
## line search fails -0.4823254 -0.5646101 1.084874e-05 7.909748e-06 -1.823444e-08 -6.944499e-09 -2.527
```

```
myCol<- rainbow(20)
myPar <- list(superpose.symbol = list(col = myCol),
              superpose.line = list(col = myCol))
ggplot(model.svmr, highlight = TRUE, par.settings = myPar)
```



```
# test error
radial_test_preds = predict(model.svmr, newdata = covid_dat[-rowTrain, ])
confusionMatrix(data = radial_test_preds ,
                 reference = covid_dat$recovery_time[-rowTrain])
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction long short
##      long    77    67
##      short  239   692
##
##              Accuracy : 0.7153
##              95% CI : (0.6873, 0.7422)
##      No Information Rate : 0.706
##      P-Value [Acc > NIR] : 0.2633
##
##              Kappa : 0.1847
##
##  Mcnemar's Test P-Value : <2e-16
##
##              Sensitivity : 0.24367
##              Specificity : 0.91173
```

```
##          Pos Pred Value : 0.53472
##          Neg Pred Value : 0.74329
##          Prevalence : 0.29395
##          Detection Rate : 0.07163
##          Detection Prevalence : 0.13395
##          Balanced Accuracy : 0.57770
##
##          'Positive' Class : long
##
```

```
#radial kernel
#set.seed(2)


#model.svmr <- tune.svm(recovery_time ~ .,
                        #data = covid_dat[rowTrain, ],
                        #kernel = "radial",
                        #cost = exp(seq(-1,5,len=20)),
                        #gamma = exp(seq(-4,3,len=20)))

#plot(model.svmr, transform.y = log, transform.x = log,
#color.palette = terrain.colors)
# test error
#radial_test_preds = predict(model.svmr, newdata = covid_dat[-rowTrain, ])
#confusionMatrix(data = radial_test_preds ,
                  #reference = covid_dat$recovery_time[-rowTrain])
```

test data performance of SVM methods

```
resamp <- resamples(list(svm1 = model.svm1,
                        svmr = model.svmr
                        ))

bwplot(resamp)
```



```
pred.svm1 <- predict(model.svm1, newdata = covid_dat[-rowTrain,])
pred.svmr <- predict(model.svmr, newdata = covid_dat[-rowTrain,])
```

```
confusionMatrix(data = pred.svm1,
                 reference = covid_dat$recovery_time[-rowTrain])
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction long short
```

```
##      long      0      0
```

```
##      short 316    759
```

```
##
```

```
##           Accuracy : 0.706
```

```
##           95% CI : (0.6778, 0.7331)
```

```
##      No Information Rate : 0.706
```

```
##      P-Value [Acc > NIR] : 0.5152
```

```
##
```

```
##           Kappa : 0
```

```
##
```

```
##      McNemar's Test P-Value : <2e-16
```

```
##
```

```
##           Sensitivity : 0.000
```

```
##           Specificity : 1.000
```

```
##      Pos Pred Value :   NaN
```

```
##      Neg Pred Value : 0.706
```

```
##           Prevalence : 0.294
```

```
##      Detection Rate : 0.000
```

```
##      Detection Prevalence : 0.000
```

```
##           Balanced Accuracy : 0.500
```

```
##
```

```
##          'Positive' Class : long
##
confusionMatrix(data = pred.svmr,
                 reference = covid_dat$recovery_time[-rowTrain])

## Confusion Matrix and Statistics
##
##          Reference
## Prediction long short
##      long      77      67
##      short 239    692
##
##          Accuracy : 0.7153
##          95% CI : (0.6873, 0.7422)
##      No Information Rate : 0.706
##      P-Value [Acc > NIR] : 0.2633
##
##          Kappa : 0.1847
##
##      McNemar's Test P-Value : <2e-16
##
##          Sensitivity : 0.24367
##          Specificity : 0.91173
##      Pos Pred Value : 0.53472
##      Neg Pred Value : 0.74329
##          Prevalence : 0.29395
##      Detection Rate : 0.07163
##      Detection Prevalence : 0.13395
##      Balanced Accuracy : 0.57770
##
##          'Positive' Class : long
##
res <- resamples(list(GLM = model.glm, GLMNET = model.glmn, GAM = model.gam, MARS = model.mars, CTREE =
trainROC <- bwplot(res, metric = "ROC")
summary(res)

##
## Call:
## summary.resamples(object = res)
##
## Models: GLM, GLMNET, GAM, MARS, CTREE, RPART, LDA, QDA, NB, SVML, SVMR
## Number of resamples: 10
##
## ROC
##          Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## GLM      0.7027786 0.7100512 0.7240580 0.7243182 0.7396824 0.7482539 0
## GLMNET 0.7010325 0.7091541 0.7199426 0.7230782 0.7374113 0.7488612 0
## GAM      0.7092686 0.7156054 0.7316766 0.7319728 0.7432725 0.7592621 0
## MARS     0.7131036 0.7220759 0.7309513 0.7309203 0.7385390 0.7542894 0
## CTREE    0.6718201 0.6806484 0.6832281 0.6925873 0.7081043 0.7224385 0
## RPART    0.6570851 0.6730605 0.6885818 0.6892499 0.6958410 0.7433196 0
## LDA      0.6982994 0.7092113 0.7204289 0.7236334 0.7385363 0.7509869 0
## QDA      0.6825470 0.6870515 0.7021897 0.7067165 0.7212097 0.7436988 0
```

```
## NB      0.6882730 0.7120171 0.7177622 0.7186405 0.7223240 0.7519739    0
## SVML    0.6968239 0.6982173 0.7103583 0.7159424 0.7324267 0.7462041    0
## SVMR    0.6404795 0.6755612 0.7070927 0.6985938 0.7242925 0.7435469    0
##
## Sens
##          Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## GLM      0.2162162 0.27027027 0.27897075 0.28146983 0.30743243 0.33783784    0
## GLMNET 0.2027027 0.24662162 0.28378378 0.26932618 0.29697334 0.32432432    0
## GAM      0.2027027 0.26277305 0.31756757 0.30170307 0.33783784 0.36486486    0
## MARS     0.2297297 0.28378378 0.29248427 0.30444280 0.33445946 0.37837838    0
## CTREE    0.1621622 0.17905405 0.19594595 0.22882266 0.28378378 0.32876712    0
## RPART    0.2027027 0.24324324 0.25000000 0.24764902 0.25939467 0.28378378    0
## LDA      0.1891892 0.26013514 0.27702703 0.26525361 0.29489078 0.31081081    0
## QDA      0.5270270 0.55743243 0.60135135 0.59546464 0.63175676 0.67567568    0
## NB       0.0000000 0.01351351 0.01351351 0.01488338 0.02369493 0.02702703    0
## SVML     0.0000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000    0
## SVMR     0.1756757 0.19451129 0.20945946 0.21917808 0.24324324 0.27027027    0
##
```

```
## Spec
##          Min.      1st Qu.      Median      Mean      3rd Qu.      Max. NA's
## GLM      0.8531073 0.9039548 0.9154605 0.9119755 0.9324890 0.9438202    0
## GLMNET 0.8644068 0.9053672 0.9180791 0.9187425 0.9382022 0.9548023    0
## GAM      0.8644068 0.8884181 0.9124294 0.9018314 0.9157303 0.9269663    0
## MARS     0.8644068 0.8912429 0.9098584 0.9057862 0.9196106 0.9325843    0
## CTREE    0.8474576 0.8997175 0.9152542 0.9103314 0.9255618 0.9604520    0
## RPART    0.8644068 0.8956151 0.9180791 0.9125627 0.9268631 0.9548023    0
## LDA      0.8587571 0.8997175 0.9239351 0.9170507 0.9324890 0.9606742    0
## QDA      0.6440678 0.7090395 0.7211325 0.7185298 0.7299562 0.7683616    0
## NB       0.9887006 0.9957627 1.0000000 0.9977401 1.0000000 1.0000000    0
## SVML     1.0000000 1.0000000 1.0000000 1.0000000 1.0000000 1.0000000    0
## SVMR     0.8870056 0.9124294 0.9378531 0.9288929 0.9423284 0.9606742    0
```

```
#pred
glm.pred <- predict(model.glm, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
glmnet.pred <- predict(model.glmn, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
gam.pred <- predict(model.gam, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
mars.pred <- predict(model.mars, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]

ctree.pred <- predict(model.ctree, newdata = covid_dat[-rowTrain,], type = "prob")[,2]
rpart.pred <- predict(model.rpart, newdata = covid_dat[-rowTrain,], type = "prob")[,2]

lda.pred <- predict(model.lda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
qda.pred <- predict(model.qda, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]
nb.pred <- predict(model.nb, newdata = covid_dat2[-rowTrain,], type = "prob")[,2]

svml.pred <- predict(model.svml, newdata = covid_dat[-rowTrain,], type = "prob")[,2]
svmr.pred <- predict(model.svmr, newdata = covid_dat[-rowTrain,], type = "prob")[,2]
```

```
#roc
roc.glm <- roc(covid_dat$recovery_time[-rowTrain], glm.pred)
roc.glmnet <- roc(covid_dat$recovery_time[-rowTrain], glmnet.pred)
roc.gam <- roc(covid_dat$recovery_time[-rowTrain], gam.pred)
roc.mars <- roc(covid_dat$recovery_time[-rowTrain], mars.pred)
```

```

roc.ctree <- roc(covid_dat$recovery_time[-rowTrain], ctree.pred)
roc.rpart <- roc(covid_dat$recovery_time[-rowTrain], rpart.pred)

roc.lda <- roc(covid_dat$recovery_time[-rowTrain], lda.pred)
roc.qda <- roc(covid_dat$recovery_time[-rowTrain], qda.pred)
roc.nb <- roc(covid_dat$recovery_time[-rowTrain], nb.pred)

roc.svml <- roc(covid_dat$recovery_time[-rowTrain], svml.pred)
roc.svmr <- roc(covid_dat$recovery_time[-rowTrain], svmr.pred)

auc <- c(roc.glm$auc[1], roc.glmn$auc[1],
         roc.gam$auc[1], roc.mars$auc[1],
         roc.lda$auc[1], roc.qda$auc[1], roc.nb$auc[1],
         roc.ctree$auc[1], roc.rpart$auc[1],
         roc.svml$auc[1], roc.svmr$auc[1]
        )

plot(roc.glm, legacy.axes = TRUE)
plot(roc.glmn, col = 2, add = TRUE)

plot(roc.gam, col = 3, add = TRUE)
plot(roc.mars, col = 4, add = TRUE)

plot(roc.lda, col = 5, add = TRUE)
plot(roc.qda, col = 6, add = TRUE)
plot(roc.nb, col = 7, add = TRUE)

plot(roc.ctree, col = 8, add = TRUE)
plot(roc.rpart, col = 9, add = TRUE)

plot(roc.svml, col = 10, add = TRUE)
plot(roc.svmr, col = 11, add = TRUE)

modelNameNames <- c("glm", "glmn", "gam", "mars", "lda", "qda", "nb", "ctree", "rpart", "svm (linear kernel)", "svm (rbf kernel)")
legend("bottomright", legend = paste0(modelNameNames, ": ", round(auc, 3)),
      col = 1:11, lwd = 2)

```

